

BY FENIL GAJJAR

KUBERNETES DAILY TASKS RBAC

(ROLE BASED ACCESS CONTROL)



COMPLETE RBAC IN KUBERNETES DOCUMENTATION

- COMPREHENSIVE GUIDE
- THEORY + PRACTICAL TASKS
- REAL TIME SCENARIO TASKS
- END TO END DOC

CONTACT US

- fenilgajjar.devops@gmail.com
- linkedin.com/in/fenil-gajjar
- github.com/Fenil-Gajjar





FOLLOW FOR MORE TASKS









Kubernetes RBAC:

Secure Your Cluster the

Right Way! 🚀



Welcome to the Ultimate Guide on **Kubernetes RBAC!**

Hey everyone!

First of all, a **BIG THANK YOU** for joining me on this exciting journey of learning Kubernetes and DevOps. 🙌

Your continuous support, kind words, and engagement truly keep me motivated to dive deeper and share more every day.

In this guide, we're going to explore one of the most crucial aspects of Kubernetes security — RBAC (Role-Based Access Control).

Whether you're a beginner trying to understand how access control works in K8s or a DevOps enthusiast building secure clusters for production — this doc is designed for you! 🚀



🔐 What This Doc Covers

Here's what you can expect to learn:

- What RBAC is and why it's essential in Kubernetes
- Key components: Roles, RoleBindings, ClusterRoles, ClusterRoleBindings
- Different kinds of subjects: Users, Groups, ServiceAccounts
- Real-time scenarios and use cases with complete setups
- How to debug and audit RBAC issues effectively
- Practical tips and best practices for secure access management

I've kept the content friendly, detailed, and hands-on — so you can understand **RBAC** deeply and apply it confidently in your clusters.

Let's Learn & Grow Together!

I'm incredibly grateful for all the support you've shown throughout my Kubernetes series. Your encouragement helps me stay consistent and keeps this momentum going strong! 🙏

So let's continue this learning journey together — step by step, one concept at a time.

Stay tuned, stay curious, and let's make Kubernetes simpler for everyone. 6



Cheers,

~ Fenil Gajjar

"Security is not a feature — it's a foundation. Let's build it right. 🔐 🚀 "

What is RBAC in Kubernetes?

RBAC (Role-Based Access Control) is a way to control who can access and perform actions on resources inside a Kubernetes cluster.

It helps you define:

- Who (user, group, or service account)
- Can do what (actions like get, create, delete, etc.)
- On which resources (like pods, deployments, secrets)
- In which namespace (or cluster-wide)

In simple terms, it helps you answer these questions:

- Who can access the cluster?
- What actions can they perform? (like create, read, update, delete)
- On which **resources**? (Pods, Deployments, Secrets, etc.)
- In which **namespace** or **scope**?

* Why is RBAC Important?

Kubernetes clusters often host **multiple applications, teams, and environments**. Without proper access control:

- A user could accidentally delete resources they shouldn't touch.
- A compromised pod might access secrets from another namespace.
- Debugging access issues becomes chaotic.

RBAC helps solve all that by enforcing the **principle of least privilege**—only granting users and services the exact permissions they need, and nothing more.



What Was Used Before RBAC in Kubernetes?

Before Kubernetes adopted the widely used RBAC (Role-Based Access Control) model, access control was managed using a more primitive method called ABAC (Attribute-Based Access Control).

While ABAC did the job in early Kubernetes days, it lacked the flexibility, scalability, and native integration needed for managing modern Kubernetes clusters securely.

What is ABAC?

ABAC is a permission model where access is granted based on a set of user-defined policies using attributes such as:

- User Identity
- Resource Type
- Namespace
- Action (Verb)

These rules are written in **JSON** files and loaded into the Kubernetes API server at startup.



```
{
  "user": "developer",
  "namespace": "dev",
  "resource": "pods",
  "verb": "create"
}
```

This allows the developer user to create pods only in the dev namespace.

↑ Why ABAC Wasn't Ideal:

Although ABAC was simple to get started with, it introduced several limitations that made it unfit for production-level systems:

- Not Dynamic: Any change to ABAC policies required restarting the Kubernetes API server — which is disruptive.
- Difficult to Manage at Scale: Large static JSON policy files became hard to maintain and audit.
- No Role Concept: ABAC didn't support reusable roles or grouping permissions — everything was user-specific.

- No Native Kubernetes Objects: Unlike RBAC, it didn't use Kubernetes-native resources like Roles and RoleBindings.
- Weaker Security Model: Fine-grained access control was difficult to achieve.

The Shift to RBAC

As Kubernetes evolved, it needed a more secure, scalable, and Kubernetes-native way to manage permissions. That's when RBAC was introduced and became generally available in Kubernetes v1.6.

From that point onward:

- RBAC became the **default and recommended** authorization mechanism.
- ABAC is still available but is **discouraged** and rarely used especially in production.



Why Do We Need RBAC in Kubernetes?

In any production-grade Kubernetes environment, security and controlled access are critical. As your cluster grows with more users, services, and teams working together — it becomes essential to control who can do what inside the cluster.

That's where **RBAC** (Role-Based Access Control) steps in. It helps you secure, manage, and organize access to Kubernetes resources in a scalable, maintainable way.

Here's Why RBAC is So Important:

1 Fine-Grained Access Control

RBAC lets you define **precise permissions**. Instead of giving users full admin access, you can allow:

- Only reading pods
- Only creating deployments
- Only viewing logs This reduces risk and enforces least privilege access.

2 Team-Based Role Assignment

With RBAC, you can assign Roles to groups or teams (like developers, DevOps, testers) rather than individual users.

- Developers → Can deploy apps in the dev namespace
- DevOps → Can manage nodes and networking
- QA → Can read logs and run test jobs
 Makes management easier and avoids permission chaos.

3 Namespace-Level Isolation

Kubernetes is multi-tenant by nature. RBAC lets you **restrict actions within namespaces**, so:

- Devs can't touch production resources
- Testers can't modify deployments
 Perfect for separating teams, stages (dev/test/prod), or applications.

4 Improves Security & Compliance

RBAC enforces who can access what, which is crucial for:

- Auditing and logging activity
- Meeting compliance standards (ISO, HIPAA, SOC2, etc.)
- Reducing the blast radius of misconfigurations or attacks

5 Kubernetes-Native & Dynamic

RBAC is built into Kubernetes:

- ullet You can **modify roles and bindings on the fly** no need to restart anything
- It's declarative you manage it like any other YAML resource
- Works seamlessly with Kubernetes tools, CLIs, and APIs

Marks How RBAC Works in Kubernetes

RBAC (Role-Based Access Control) in Kubernetes allows you to control access to resources within the cluster by defining **Roles** and **RoleBindings** (or **ClusterRoles** and **ClusterRoleBindings** for cluster-wide access).

1 Roles & ClusterRoles

At the heart of RBAC are the **Roles** and **ClusterRoles**. These define the **permissions** (what actions can be performed) for a specific set of resources.

- Role: This defines permissions within a namespace. For example, you
 might want to allow users to create and read pods within a specific
 namespace but not in the entire cluster.
- ClusterRole: A ClusterRole is similar to a Role but can apply across the
 entire cluster, including all namespaces. It's used for cluster-wide
 resources like nodes or persistent volumes.

2 RoleBindings & ClusterRoleBindings

Once you have **Roles** or **ClusterRoles**, you then need to bind them to users, groups, or service accounts using **RoleBindings** and **ClusterRoleBindings**. These bindings grant the roles to a specific **user** or **service account**, giving them the defined permissions.

 RoleBinding: A RoleBinding grants the permissions of a Role within a specific namespace to a user or set of users. ClusterRoleBinding: A ClusterRoleBinding grants the permissions of a
 ClusterRole across the entire cluster.

This way, you can create flexible access control policies, whether you want to restrict actions to a particular namespace or give global access to cluster-wide resources.

3 Kubernetes Resources & Permissions

Roles define the types of actions that can be performed on **Kubernetes resources**. These resources could be things like **pods**, **deployments**, **services**, **configmaps**, etc. Permissions typically fall into **verbs** like:

- get (view)
- list (list multiple resources)
- create (create resources)
- update (update resources)
- delete (remove resources)

For example, a Role may allow a user to get and list pods but not create or delete them.

4 Service Accounts & RBAC

In addition to users, **service accounts** can be assigned roles. Service accounts are used by **pods** to access the Kubernetes API. This allows you to control and restrict **what the pod can do**. For example, you might have a pod that only needs to read from a specific set of resources and not modify them.

5 RBAC Workflow: Step-by-Step

Here's how RBAC works in Kubernetes:

- Define Roles (or ClusterRoles): Specify what actions can be performed on which resources.
 - Example: A Role allowing create, list, and get permissions on pods in a specific namespace.
- Create RoleBindings (or ClusterRoleBindings): Bind the defined Roles or ClusterRoles to specific users, groups, or service accounts.
 - Example: A RoleBinding assigns the Role created in step 1 to a developer, allowing them to create and manage pods in the development namespace.
- Access Control: Now, when the user or service account makes requests to Kubernetes, RBAC determines whether they are allowed to perform the requested actions based on the Roles and Bindings.

6 RBAC Example in Action

Let's break down a simple example to understand how everything ties together:

Step 1: You define a **Role** (named dev-namespace-admin) that allows users to create, get, list, and delete pods in the dev namespace:

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
   namespace: dev
   name: dev-namespace-admin
rules:
- apiGroups: [""]
   resources: ["pods"]
   verbs: ["create", "get", "list", "delete"]
```

Step 2: You then create a **RoleBinding** that binds this role to a specific user (developer-user) within the dev namespace:

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
```

name: dev-namespace-admin-binding

namespace: dev

subjects:

- kind: User

name: developer-user

apiGroup: rbac.authorization.k8s.io

roleRef:

kind: Role

name: dev-namespace-admin

apiGroup: rbac.authorization.k8s.io

• Step 3: The developer-user is now able to perform the specified actions (create, get, list, delete) on pods in the dev namespace. If they attempt to access resources outside this namespace or try other actions (like update), they will be denied.



RBAC (Role-Based Access Control) in Kubernetes provides a **powerful and flexible** way to manage access control within your cluster. It allows you to define **who can access what**, and **what actions** they can perform on specific resources.
Let's explore the essential components that make up RBAC!

X 1. Role

A **Role** defines a set of permissions (like **get**, **create**, **list**, **delete**) for resources within a **specific namespace**. It is ideal for managing access to resources like **Pods**, **Deployments**, **ConfigMaps**, and more within that namespace.

Example Use Case:

• A **Role** could allow a user to **list** and **get** Pods in the **dev** namespace but **not** allow them to create or delete Pods.

2. ClusterRole

A **ClusterRole** is like a **Role**, but with permissions that apply **cluster-wide** (not just within a namespace). ClusterRoles are useful for managing access to resources that span across the entire Kubernetes cluster, such as **Nodes** or **PersistentVolumes**.

Example Use Case:

 A ClusterRole could allow a user to manage all nodes in the cluster or view metrics across multiple namespaces.

A **RoleBinding** connects a **Role** to a specific **user**, **group**, or **service account** within a **namespace**. It ensures that the permissions granted by the Role apply only within that namespace.

Example Use Case:

• A **RoleBinding** can assign the **developer-role** to the **developer-user** in the **dev** namespace, enabling them to perform actions like **get** and **list** Pods.

4. ClusterRoleBinding

A ClusterRoleBinding connects a ClusterRole to a user, group, or service account, but applies across the entire cluster. This gives the subject access to resources and actions across all namespaces and cluster-wide resources.

Example Use Case:

 A ClusterRoleBinding can assign the admin-cluster-role to a superuser, granting full access to resources like Nodes and Persistent Volumes.



Subjects are the **users**, **groups**, or **service accounts** that are granted permissions through **RoleBindings** or **ClusterRoleBindings**. These are the entities that will have access to perform actions defined in the **Role** or **ClusterRole**.

Example Use Case:

A subject could be a specific user (developer-user), a group (dev-team),
 or a service account (dev-service-account).

📌 6. RoleRef

The **RoleRef** in a **RoleBinding** or **ClusterRoleBinding** references the **Role** or **ClusterRole** that defines the permissions being granted. It's the link between the subject and the set of permissions they will have access to.

Example Use Case:

• A **RoleBinding** could reference a **developer-role** under **roleRef**, ensuring the assigned user has the correct permissions.

☆ 7. Verbs

Verbs define the **actions** that can be performed on resources. These actions are what you specify when creating a **Role** or **ClusterRole** to control what operations a subject can perform on specific resources.

Common Verbs include:

get: Retrieve resource information

- list: List multiple resources
- create: Create a new resource
- update: Modify an existing resource
- delete: Remove a resource
- patch: Partially update a resource
- watch: Monitor resources for changes

8. Resources

Resources refer to the **types of objects** in Kubernetes that RBAC manages access to. Some examples of resources are:

- Pods
- Deployments
- Services
- ConfigMaps
- Secrets
- Nodes (for cluster-wide access)

These resources can have specific permissions associated with them, enabling fine-grained access control.



🔐 Deep Dive into Roles and RoleBindings in Kubernetes 🚀



In Kubernetes, Role and RoleBinding are core components of Role-Based Access Control (RBAC) that manage who can access what resources within a Kubernetes cluster, and what actions they can perform on those resources. Understanding how to configure **Roles** and **RoleBindings** is essential for secure and efficient management of your Kubernetes environment.

What is a Role?

A Role defines a set of permissions (known as verbs) that can be applied to resources within a specific namespace. The role allows administrators to control what actions users and services can take on resources within that namespace. A Role is **namespace-specific**, meaning it is only effective within the namespace where it is created.

Key Elements of a Role

- 1. **Verbs**: These define the actions a subject can perform on the resource. Common verbs include:
 - get: Retrieve a resource.
 - list: List multiple resources.
 - create: Create a resource.

0	delete: Remove a resource.
0	update: Modify a resource.
0	patch: Apply a partial update to a resource.
0	watch: Monitor resources for changes.
Resources : These are the objects or entities within the cluster that users can interact with. Examples include:	
0	pods
0	services
0	deployments
0	configmaps
0	secrets
0	nodes (if used in a ClusterRole)
API Groups : Resources in Kubernetes belong to different API groups. For example, apps for deployments, core for pods, and so on. The Role defines permissions for resources in a specific API group.	

2.

3.

4. **Namespaces**: Roles are always bound to a **specific namespace** (except for ClusterRoles, which apply to the entire cluster). This means that a Role only governs permissions for resources within that namespace.

What is a RoleBinding?

A **RoleBinding** is used to **assign a Role** to a **user**, **group**, or **service account**. It defines which users or services should have the permissions defined in the **Role**. When a **RoleBinding** is created, the permissions specified in the Role are applied to the subject (user/group/service account).

A **RoleBinding** works within a **specific namespace**, meaning it is bound to the namespace where it is created. It ensures that the subject has access to the resources defined in the Role within that namespace.

Key Elements of a RoleBinding

RoleRef:

 This references the Role or ClusterRole that defines the set of permissions being granted. The RoleRef essentially links the Role to the user or service account being granted access.

2. Subjects:

 The subjects are the entities (users, groups, or service accounts) to which the Role is granted. This is where the permissions from the Role are applied.

How Role and RoleBinding Work Together

1. Create the Role:

 First, a Role is created in a specific namespace, detailing the permissions (verbs) and resources (like pods, services) that a user or service can interact with.

2. Bind the Role to Users:

 After the Role is created, a **RoleBinding** is created to bind that Role to a user or service account. The **RoleBinding** grants the user the permissions defined in the Role.

3. Access Control:

 When a user or service account tries to access a resource in the namespace, the RoleBinding checks if they have the correct permissions. If the user has the necessary role-bound permissions, they are granted access; otherwise, access is denied.

Example of a Role and RoleBinding Configuration

Here's an example of creating a **Role** and **RoleBinding** to allow a user to **view** (but not modify) Pods in the **dev** namespace:

1 Role Definition:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  # Name of the Role
  name: pod-viewer
  # Specify the namespace where the role is applicable
  namespace: dev
rules:
  - verbs:
      - get
      - list
    resources:
      - pods
```

In this Role:

- The user can **get** and **list** Pods within the **dev** namespace.
- The Role doesn't allow the user to **create**, **delete**, or **update** Pods.

2 RoleBinding Definition:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  # Name of the RoleBinding
  name: pod-viewer-binding
  # Specify the namespace of the RoleBinding
  namespace: dev
subjects:
  - kind: User
    name: developer-user
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
```

name: pod-viewer

apiGroup: rbac.authorization.k8s.io

In this RoleBinding:

• The **developer-user** is granted the **pod-viewer** Role.

 The permissions to get and list Pods in the dev namespace are assigned to the user.

* Why Are Roles and RoleBindings Important?

1. Fine-Grained Access Control:

 Roles give you the flexibility to define precise permissions for users, ensuring that they only have access to the resources they need. By binding roles to specific namespaces, you ensure that users can only access the right resources.

2. Least Privilege Principle:

 Roles and RoleBindings help enforce the least privilege principle, ensuring that users and service accounts only have access to the minimum resources necessary to do their jobs.

3. Security:

 By using Roles and RoleBindings, you can limit the scope of access to Kubernetes resources, preventing unauthorized access or accidental misuse of sensitive resources like secrets and deployments.

4. Namespace Isolation:

 Roles and RoleBindings support Kubernetes' namespace isolation, which helps segregate permissions across multiple environments, such as development, staging, and production.

Best Practices for Using Roles and RoleBindings

1. Create Specific Roles for Specific Tasks:

 Avoid giving broad permissions (e.g., admin) unless absolutely necessary. Instead, create roles tailored to specific tasks.

2. Use RoleBindings for Temporary Access:

 If a user only needs temporary access to certain resources, bind them to a role with the required permissions for a limited period.

3. Avoid Overusing ClusterRoles in RoleBindings:

ClusterRoles are more general and apply across the entire cluster.
 Use them when needed, but try to restrict roles to specific namespaces whenever possible.

4. Audit and Review Regularly:

 Regularly review the Roles and RoleBindings in your cluster to ensure they are still appropriate. Ensure no unnecessary permissions are granted, and revoke access when users no longer need it.

What Happens If a RoleBinding Refers to a Role in a Different Namespace?

In Kubernetes, RoleBinding is namespace-specific and can only bind to Roles within the same namespace. If you try to create a RoleBinding that references a Role in a different namespace, it will not work and will result in an error. This is because Kubernetes strictly enforces the namespace boundary for Roles and RoleBindings.

Why Is This the Case?

- Roles are namespace-scoped, meaning they define permissions for resources only within the namespace they are created in.
- RoleBindings, by design, must be tied to the same namespace as the Role
 they reference. The RoleBinding gives the subject (user/service account)
 the permissions granted by the Role within that same namespace.

Kubernetes enforces this rule to maintain clear **namespace isolation**, ensuring that resources and permissions are managed independently in each namespace.

💡 What Happens If You Do It Anyway?

If you attempt to create a **RoleBinding** that references a **Role** in a different namespace, Kubernetes will reject the request with an error like:

error: roles.rbac.authorization.k8s.io "role-name" is forbidden: role "role-name" is in a different namespace

This error indicates that the Role and RoleBinding are not compatible because they exist in different namespaces.

✓ Correct Way to Handle Cross-Namespace Access:

If you need to **grant access to resources across multiple namespaces**, you can use **ClusterRoles** and **ClusterRoleBindings**.

- ClusterRole: Unlike Roles, a ClusterRole is cluster-wide and can be used across multiple namespaces. It defines permissions for resources across the entire cluster, not bound to any particular namespace.
- ClusterRoleBinding: If you want to grant a user access to resources across different namespaces (or even the entire cluster), you should use a ClusterRoleBinding instead of a RoleBinding.

Example of Correct Usage:

1 ClusterRole Definition (Cluster-wide Permissions):

apiVersion: rbac.authorization.k8s.io/v1

```
kind: ClusterRole
metadata:
  name: pod-reader
rules:
  - verbs:
       - get
       - list
     resources:
       - pods
2 ClusterRoleBinding Definition (Binding Across Namespaces):
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: pod-reader-binding
subjects:
  - kind: User
    name: "developer-user"
    apiGroup: rbac.authorization.k8s.io
```

roleRef:

kind: ClusterRole

name: pod-reader

apiGroup: rbac.authorization.k8s.io

In this example:

- The ClusterRole pod-reader grants permission to view Pods cluster-wide.
- The **ClusterRoleBinding** binds the **developer-user** to this ClusterRole, allowing them to view Pods across any namespace in the cluster.

® Different Kinds of RoleBinding in Kubernetes **®**

In Kubernetes, a **RoleBinding** is used to assign **roles** to **users**, **groups**, or **service accounts** within a specific **namespace**. RoleBinding essentially grants the permissions defined in a **Role** or **ClusterRole** to a set of subjects (users, groups, or service accounts).

While **RoleBinding** itself is quite specific to the namespace, there are a few different **subjects** and **role references** you can use to customize how you assign roles.

🔳 RoleBinding with Role 🎭

This is the most common **RoleBinding** used in Kubernetes. It links a **Role** (which is namespace-scoped) to subjects such as users, service accounts, or groups.

Example:

Let's say you want to allow a **user** to list and get Pods in a specific namespace.

Role (within namespace dev):

apiVersion: rbac.authorization.k8s.io/v1

kind: Role

metadata:

namespace: dev

name: pod-reader

rules:

```
- verbs:
      - get
      - list
    resources:
      - pods
RoleBinding:
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: pod-reader-binding
  namespace: dev
subjects:
  - kind: User
    name: "dev-user"
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
```

apiGroup: rbac.authorization.k8s.io

In this example:

The Role pod-reader in the dev namespace grants permission to get and

list Pods.

The **RoleBinding** binds this Role to the **user** dev-user in the dev

namespace.

2 RoleBinding with ClusterRole 🌍

In certain situations, you may want to grant cluster-wide permissions but within a

specific namespace. For this, you can bind a ClusterRole to a subject using

RoleBinding. This allows a user or service account to have permissions granted

by a **ClusterRole** but confined to the namespace where the **RoleBinding** exists.

Example:

Let's say you want to allow a user to access Pods across the cluster, but in the dev

namespace only. You can use a **ClusterRole** for cluster-wide permissions and bind

it in a namespace.

ClusterRole:

apiVersion: rbac.authorization.k8s.io/v1

kind: ClusterRole

```
metadata:
  name: pod-reader-cluster
rules:
  - verbs:
      - get
      - list
    resources:
      - pods
RoleBinding (Namespace-specific):
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: pod-reader-binding
  namespace: dev
subjects:
  - kind: User
    name: "dev-user"
    apiGroup: rbac.authorization.k8s.io
```

roleRef:

kind: ClusterRole

name: pod-reader-cluster

apiGroup: rbac.authorization.k8s.io

In this case:

- The ClusterRole pod-reader-cluster grants permissions to get and list Pods across the entire cluster.
- The **RoleBinding** in the dev namespace grants dev-user the cluster-wide permissions within that namespace.

3 RoleBinding with ServiceAccount

A **ServiceAccount** is often used in Kubernetes to allow applications to interact with the cluster. You can use a **RoleBinding** to assign roles to a **ServiceAccount**.

Example:

If you want to allow a **ServiceAccount** to get and list Pods in the dev namespace, you could do something like this:

Role (within namespace dev):

apiVersion: rbac.authorization.k8s.io/v1

```
kind: Role
metadata:
  namespace: dev
  name: pod-reader
rules:
  - verbs:
    - get
    - list
  resources:
    - pods
```

RoleBinding for ServiceAccount:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
   name: pod-reader-binding
   namespace: dev
subjects:
```

- kind: ServiceAccount

name: "dev-service-account"

namespace: dev

roleRef:

kind: Role

name: pod-reader

apiGroup: rbac.authorization.k8s.io

In this example:

• The **Role** grants permissions to get and list Pods in the dev namespace.

The RoleBinding binds this Role to a ServiceAccount

dev-service-account in the dev namespace.

4 RoleBinding with Group 👥

In Kubernetes, you can also bind a **Role** to a group of users. Groups are not directly created in Kubernetes, but they can be defined externally (e.g., through your identity provider or Kubernetes OIDC configuration).

Example:

If you want to allow a specific **group** of users to list Pods in the dev namespace:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: dev
  name: pod-reader
rules:
  - verbs:
      - get
      - list
    resources:
      - pods
RoleBinding for Group:
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: pod-reader-binding
  namespace: dev
```

Role (within namespace dev):

subjects:

```
- kind: Group
  name: "dev-group"
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

In this example:

- The **Role** grants permissions to get and list Pods in the dev namespace.
- The **RoleBinding** binds this **Role** to the **group** dev-group.

What is kind: User in Kubernetes RBAC?

When you define a subject in a RoleBinding or ClusterRoleBinding, you're specifying **who** should get the permissions. That who is called a **subject**, and it could be one of the following:

- User
- Group
- ServiceAccount

So in this example:

- kind: User

name: "dev-user"

You're saying:

"I want to give access to a user named dev-user."

But what exactly is a "User" in Kubernetes?

Kubernetes **does not** manage user accounts by itself like traditional systems (e.g., Linux). Instead, **users are managed externally** – such as via:

- Certificates (for example, kubect1 uses client certs to authenticate)
- V Identity providers via OIDC (like Google, Azure AD, Okta)

- V Authentication plugins
- Static username/password (for test setups, not recommended for production)

Examples of Kubernetes "User" identities:

- A user authenticated with a client certificate:
 e.g. a kubectl user configured with .kube/config using certs.
- A user authenticated via OIDC provider:
 e.g. fenil.gajjar@example.com from your organization's SSO system.
- A static username configured in an **authentication webhook**.

@ Goal

• You have an IAM user (e.g.,

```
arn:aws:iam::111122223333:user/dev-user)
```

- You want to grant them limited permissions in EKS, such as:
 - $\circ\ \ \,$ Only read access to name spaces, pods, and deployments
- So you'll:

- Map the IAM user to a custom Kubernetes group in the aws-auth
 ConfigMap
- Then create a Role + RoleBinding (or ClusterRole + ClusterRoleBinding) in Kubernetes for that group



1 IAM User Mapping in aws-auth ConfigMap

Edit the aws-auth ConfigMap:

kubectl edit configmap aws-auth -n kube-system

Add the user under mapUsers like this:

```
mapUsers: |
```

- userarn: arn:aws:iam::111122223333:user/dev-user

username: dev-user

groups:

- eks-readonly-group



- We're assigning the IAM user to a custom group eks-readonly-group.
- This group name is arbitrary; just make sure it matches in your RBAC definition.

2 Create a Kubernetes ClusterRole

This defines what permissions the group has. Example: Read-only access to core resources.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
    name: eks-readonly-role
rules:
    - apiGroups: [""]
    resources: ["pods", "services", "namespaces"]
    verbs: ["get", "list", "watch"]
- apiGroups: ["apps"]
    resources: ["deployments", "replicasets"]
    verbs: ["get", "list", "watch"]
```

3 Bind the Role to the Group

Now bind the above ClusterRole to the group eks-readonly-group.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: eks-readonly-binding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: eks-readonly-role
subjects:
  - kind: Group
    name: eks-readonly-group
    apiGroup: rbac.authorization.k8s.io
```

Now, any IAM user in that group (like dev-user) can only view pods, deployments, namespaces, etc., and cannot create/update/delete anything.

Test the IAM User

Log in as the IAM user, set credentials, then:

aws eks update-kubeconfig --region <region> --name <cluster-name>

kubectl get pods -A

kubectl delete pod <pod-name> # X This should fail

Rest Practice

Action Recommended? X Avoid unless necessary Give system:masters Yes, for least privilege Use custom RBAC groups Use ClusterRole vs ClusterRole for global access across Role

namespaces

Real-world Scenario

Your DevOps team has multiple IAM users:

• dev-user-1, dev-user-2, qa-user

You assign all of them to eks-readonly-group in aws-auth and bind this group to a ClusterRole that provides read-only access.

This allows them to **monitor resources** without the risk of breaking things.

You don't create the group in Kubernetes directly.

In EKS, when you map an **IAM user or role** in aws-auth ConfigMap and assign a group (like eks-readonly-group), that group **doesn't need to exist** as a Kubernetes object. It's not like a Linux group or a K8s resource you create.

The group is just a **label or string** that Kubernetes uses to match in RBAC bindings.

Let's explain this in detail 👇

Where Does kind: Group Come From?

In Kubernetes RBAC:

- A Group is a logical identity.
- It's not something you manually create.
- It's used to match subjects in RoleBinding or ClusterRoleBinding.

So, when you define this:

subjects:

- kind: Group

name: eks-readonly-group

Kubernetes checks:

"Is the user (or role) calling me associated with this group?"

The answer depends on how the IAM identity was mapped in aws-auth.

Mapping in aws-auth Makes It Work

When you do this:

```
mapUsers: |
```

- userarn: arn:aws:iam::111122223333:user/dev-user

username: dev-user

groups:

- eks-readonly-group

You're telling Kubernetes:

"When this IAM user calls kubectl, treat them as a user named dev-user and a member of the group eks-readonly-group."

Now, if your RBAC bindings say "grant permissions to eks-readonly-group", it matches and applies access.

X How This Works in Practice

Here's the full interaction:

- 1. IAM user dev-user runs kubectl get pods
- 2. kubectl signs the request with AWS credentials
- 3. EKS authenticates the request using IAM
- 4. It uses aws-auth to translate the IAM identity:
 - o Maps to username: dev-user
 - Maps to group: eks-readonly-group
- Kubernetes checks all RoleBindings/ClusterRoleBindings for eks-readonly-group
- 6. If a match is found, it allows/denies access accordingly

No Need to Create the Group

You don't need to run kubectl create group or anything like that.

It's a virtual identity label.

Just make sure:

- The group name in aws-auth matches
- The group name in RoleBinding/ClusterRoleBinding matches

Then you're good!

Want to See This in Action?

Run this to see the user's identity and group used when making requests:

kubectl auth can-i get pods --as=dev-user
--as-group=eks-readonly-group

This simulates a request **as that user + group**, so you can verify your RBAC.

What You Need to Configure for the IAM User

Let's say the IAM user is dev-user.

1 Ensure IAM User Has eks: DescribeCluster Permission

This is **required** to run aws eks update-kubeconfig — it fetches cluster endpoint & certificate.

Add a policy to the IAM user like this:

You can also give broader EKS permissions depending on the role.

2 Install AWS CLI and Configure Credentials

On the IAM user's machine (or DevOps environment):

```
aws configure
```

Enter:

- Access Key
- Secret Key
- Region
- Output format (e.g., json)

Or use environment variables:

```
export AWS_ACCESS_KEY_ID=xxxx
export AWS_SECRET_ACCESS_KEY=xxxx
export AWS_DEFAULT_REGION=us-east-1
```

3 Update kubeconfig for the EKS Cluster

This is a **must** so the IAM user can interact with the cluster using kubectl.

aws eks update-kubeconfig --region <region> --name
<cluster-name>

This command will:

- Add/update a context in ~/.kube/config
- Set up the cluster endpoint, CA cert, and authentication

4 Test with kubectl

Once kubeconfig is updated, try this:

kubectl get pods --all-namespaces

- If the IAM user was added to a custom group like eks-readonly-group, and RBAC allows read access they will succeed.
- X If they try something not allowed (e.g., kubectl delete pod), RBAC will deny it.

When Should You Configure the IAM User?

You're absolutely right:

You only need to configure AWS credentials & run aws eks update-kubeconfig if the IAM user wants to interact with the EKS cluster using kubect1.

What Happens If You Don't Configure the IAM User?

If you've done the following:

- Added IAM user to aws-auth ConfigMap
- Mapped to proper RBAC group
- Created appropriate Role/RoleBinding or ClusterRoleBinding

...but you **don't configure** aws configure or update-kubeconfig on their system — then:

X That IAM user won't be able to interact with the cluster — because they don't have any way to **authenticate** and **access the cluster endpoint**.

RBAC alone is not enough. RBAC is for authorization, not authentication.

What Happens When You Run aws eks

update-kubeconfig?

This command **modifies** (or creates) the ~/.kube/config file to allow your **IAM** user to interact with the EKS cluster.

It adds:

- 1. A new cluster entry
- 2. A new user entry
- 3. A new context that links the two

Let's Look at a Real Example

After running the command, your kubeconfig may look like this:

```
apiVersion: v1
kind: Config
clusters:
- cluster:
    server: https://<EKS-endpoint>
    certificate-authority-data: LS0tLS1... # base64 encoded
CA cert
```

```
name:
arn:aws:eks:us-east-1:111122223333:cluster/dev-cluster
users:
- name:
arn:aws:eks:us-east-1:111122223333:cluster/dev-cluster
  user:
    exec:
      apiVersion: "client.authentication.k8s.io/v1beta1"
      command: "aws"
      args:
        - "eks"
        - "get-token"
        - "--cluster-name"
        - "dev-cluster"
      # This uses your current AWS IAM identity to fetch a
temporary token
contexts:
- context:
```

cluster:

arn:aws:eks:us-east-1:111122223333:cluster/dev-cluster

user:

arn:aws:eks:us-east-1:111122223333:cluster/dev-cluster

name:

arn:aws:eks:us-east-1:111122223333:cluster/dev-cluster

current-context:

arn:aws:eks:us-east-1:111122223333:cluster/dev-cluster

What Each Section Means

Section	Purpose
clusters	Contains cluster API endpoint and CA cert
users	Defines how you authenticate — in this case, using aws eks get-token
contexts	Combines a cluster and user
current-co ntext	The default context used when you run kubect1

Authentication via aws eks get-token

- It uses your IAM credentials (aws configure) to generate a temporary token.
- The token is passed to the Kubernetes API server as a Bearer token.
- The cluster uses the aws-auth ConfigMap to map this IAM identity to a Kubernetes username/group.
- RBAC rules apply based on that mapping.

** So after configuring, what changes?

Before	After update-kubeconfig
No access to cluster	kubeconfig knows how to reach the cluster
No authentication	kubeconfig uses AWS CLI to generate token using IAM
No context	Now you can switch, use, and interact via kubect1

Bonus Tip: View the Config

To view your current kubeconfig:

kubectl config view --minify

To list all contexts:

kubectl config get-contexts

To switch:

kubectl config use-context <context-name>

Does current-context Change When I Run aws eks update-kubeconfig?

Yes, by default it will **set the current context** to the new EKS cluster you're configuring.

So if you're using:

aws eks update-kubeconfig --region us-east-1 --name dev-cluster

It adds the new context and also sets:

current-context:

arn:aws:eks:us-east-1:xxxxxxxxxxxxxccluster/dev-cluster

So yes — if this IAM user is used to configure kubeconfig on a machine, the current-context becomes theirs by default.

Can kubeconfig hold multiple contexts?

Yes! The kubeconfig file can hold multiple clusters, users, and contexts.

But...

1 Only one context can be "current" at a time

The current-context is what kubectl uses when you don't specify --context.

You can view all contexts with:

kubectl config get-contexts

You can switch to another one using:

kubectl config use-context <context-name>

Example: Multiple Users or Clusters in One Kubeconfig

contexts:

- name: admin-context

context:

cluster: dev-cluster

user: admin-user

- name: readonly-context

context:

cluster: dev-cluster

user: readonly-user

current-context: admin-context

At any time, only one of these will be active (current).

So, in your scenario:

- You mapped IAM user to RBAC group
- You configured it with aws eks update-kubeconfig
- Now this IAM user's context becomes current
- But you can still switch back to your previous context if needed

Here are the most **realistic and practical scenarios** where this is implemented:

✓ 1. Secure Access Control for Teams

★ Scenario:

You have multiple team members (Dev, QA, DevOps) who need access to the same EKS cluster, but with **different permissions**.

Solution:

- Map each IAM user or IAM role to a Kubernetes group via aws-auth
- Use **RBAC RoleBindings** to give scoped access

Example:

IAM User	Mapped	RBAC Role
	Group	
dev_user	eks-dev-gr oup	Can only get, list, watch pods in dev namespace
qa_user	eks-qa-gro up	Can only read ConfigMaps & secrets in qa namespace

admin_us eks-admins Full access (ClusterRoleBinding to er cluster-admin)

2. Access Control for CI/CD Pipelines (IAM Roles)

★ Scenario:

You're running **GitHub Actions**, **Jenkins**, or **GitLab CI** — and need your pipeline to deploy apps to EKS securely.

③ Solution:

- Use IAM Role for your CI/CD pipeline with sts:AssumeRole or OIDC
- Map that IAM Role in aws-auth
- Bind it to an RBAC ClusterRole with limited permissions (e.g., only in prod namespace)

☑ 3. Granular Namespace Isolation in Multi-Tenant EKS

★ Scenario:

You run a **multi-tenant platform** (e.g., SaaS), where each customer has a namespace in the same EKS cluster.

③ Solution:

- Create a unique IAM Role/User per tenant
- Map them in aws-auth and to Kubernetes group
- Apply strict RBAC bindings to their namespace only
- Prevent cross-namespace access

4. Auditing and Least Privilege Enforcement

★ Scenario:

Your security team wants **accountability and traceability** — they want to know who did what in the EKS cluster.

© Solution:

- Use individual IAM users or per-team IAM roles
- Map those identities to unique Kubernetes users/groups
- Use CloudTrail for IAM actions + EKS audit logs for RBAC

5. Federated Access (SSO Integration)

★ Scenario:

Your organization uses **AWS SSO**, **Okta**, or **Azure AD**, and developers log in via federated identities.

© Solution:

- Federated login gives IAM Role via SAML/OIDC
- That IAM Role is mapped to EKS groups via aws-auth
- RBAC defines what they can do once logged in

✓ 6. External Dev Access in Dev/Test Clusters

★ Scenario:

You give external consultants or interns temporary access to your dev EKS cluster.

⊚ Solution:

- Create IAM user with limited access
- Map to a low-privilege RBAC role (view-only)
- Rotate/remove when no longer needed

EKS vs Kops — Authentication & Access Control

Differences

Feature	EKS	Kops
Authenticati on	IAM-based + aws-auth config	Mostly kubeconfig with user/pass or token
RBAC	Based on IAM → Group → RoleBinding	Based on Kubernetes-native ServiceAccount, user certs, or tokens
Best for	AWS-native integrations, federated access, CI/CD IAM roles	Custom Kubernetes clusters, self-managed auth, more flexibility

✓ In Kops — You Typically Use:

1. ServiceAccount + RBAC

This is best for apps/pipelines/pods needing in-cluster access.

- Create a ServiceAccount
- Bind it via Role or ClusterRoleBinding
- Generate token manually or via automation

- Use it in kubeconfig file
- Common for CI/CD, monitoring agents, etc.

2. User Certificate Auth (for humans)

If a human user wants access:

- Generate client cert + private key
- Create a kubeconfig using that cert/key
- Use RBAC to bind to roles
- Full control but less dynamic than IAM

3. Token Auth (Static or Dynamic)

- You can generate long-lived tokens for users manually
- Or use **OIDC/dex** for dynamic auth

"This IAM user/role mapping is useful in EKS, but in Kops, for scenarios like developer or pipeline access, I should use ServiceAccount + RBAC + token + kubeconfig."



What is a ClusterRole in Kubernetes?

A ClusterRole is like a Role, but it is not limited to a single namespace. It can grant permissions:

- Across the entire cluster
- Or for **non-namespaced resources** like nodes, persistent volumes, etc.
 - ★ Use ClusterRole when:
 - You need to give access across multiple namespaces
 - You're dealing with cluster-wide resources

What is a ClusterRoleBinding?

A ClusterRoleBinding connects:

• A ClusterRole → to → a user, group, or service account

It tells Kubernetes:

This identity should have these cluster-wide permissions."

Example: Grant Cluster-Wide Read Access to All Resources

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: read-only
rules:
  - apiGroups: [""]
    resources: ["pods", "services", "configmaps"]
    verbs: ["get", "list", "watch"]
  - apiGroups: ["apps"]
    resources: ["deployments", "statefulsets"]
    verbs: ["get", "list", "watch"]
    This ClusterRole gives read-only access to core and apps
    resources.
```

⊗ ClusterRoleBinding Example (for a User)

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
```

metadata: name: read-access-binding subjects: - kind: User name: dev-user apiGroup: rbac.authorization.k8s.io roleRef: kind: ClusterRole name: read-only apiGroup: rbac.authorization.k8s.io

This binds the read-only ClusterRole to a user named dev-user.

Real-World Example: Give EKS IAM Group Read Access

If you mapped a group in aws-auth like this:

groups:

- eks:readonly

You can bind it like this:

subjects:

- kind: Group

name: eks:readonly

apiGroup: rbac.authorization.k8s.io

Different Kinds of subjects.kind in RBAC:

Here are the valid values for kind:

1. User 👤

Represents an individual user identity.

In cloud setups like EKS, this often maps to an ${\bf IAM}$ user.

- kind: User

name: "dev-user"

* Example use case: Giving a specific developer read access to resources.

2. Group 👥

Represents a collection of users.

If a user belongs to a group, they inherit the group's permissions.

- kind: Group

name: "developers"

*Example use case: Granting access to all users in a group like system:masters or an IAM group in EKS.

3. ServiceAccount 🤖

Represents a **Kubernetes service account**, typically used by pods to interact with the API server.

- kind: ServiceAccount

name: app-sa

namespace: my-app

📌 Example use case: Allowing an app pod to read ConfigMaps or Secrets.

Scenario:

You have a user named dev-user (maybe an IAM user or a certificate-based user) who needs **read-only access** to Pods across **all namespaces**. This user should not be able to modify or delete anything — just **view** Pods.

1 Create a ClusterRole with read-only access to Pods

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
   name: read-pods
rules:
   - apiGroups: [""]
   resources: ["pods"]
   verbs: ["get", "list", "watch"]
```

This ClusterRole gives permission to view all Pods in any namespace.

2 Create a ClusterRoleBinding to bind the ClusterRole to the user

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: read-pods-binding
subjects:
  - kind: User
                    # Must match the user identity
    name: dev-user
in the cluster auth system
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: read-pods
  apiGroup: rbac.authorization.k8s.io
```

Note: The name: dev-user should match the user authenticated via certificate, IAM (in EKS), or your auth provider.

W

Optional Namespace-Specific Setup (Using Role &

RoleBinding)

If instead, you want to give the same read access **only in a specific namespace** (say dev), use a **Role** and **RoleBinding** instead:

Role (namespace-specific):

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
   namespace: dev
   name: read-pods
rules:
   - apiGroups: [""]
   resources: ["pods"]
   verbs: ["get", "list", "watch"]
```

RoleBinding (namespace-specific):

```
apiVersion: rbac.authorization.k8s.io/v1
```

kind: RoleBinding

metadata:

namespace: dev

name: read-pods-binding

subjects:

- kind: User

name: dev-user

apiGroup: rbac.authorization.k8s.io

roleRef:

kind: Role

name: read-pods

apiGroup: rbac.authorization.k8s.io

Name :

You have a group of developers (dev-team) who should have **read-only access** to all resources in the dev namespace. Instead of assigning roles user by user, you want to assign permissions to the **entire group**.

***** Step-by-step Setup Using Group-Based RBAC

1. Create a Role (namespace-scoped)

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: dev-team-read-access
  namespace: dev
rules:
  - apiGroups: [""]
  resources: ["pods", "services", "configmaps"]
  verbs: ["get", "list", "watch"]
```

This role gives get, list, and watch access to Pods, Services, and ConfigMaps in the dev namespace.

• 2. Create a RoleBinding for the Group

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: dev-team-rolebinding
  namespace: dev
subjects:
  - kind: Group
                                        # Group name from
    name: dev-team
IAM/OIDC/LDAP etc.
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: dev-team-read-access
  apiGroup: rbac.authorization.k8s.io
```

This binds the dev-team group to the Role we just created in the dev namespace.

Want to grant cluster-wide access instead?

Let's say the **same dev-team group** should be able to view **Nodes and PersistentVolumes** cluster-wide. For that, use ClusterRole and
ClusterRoleBinding.

3. Create a ClusterRole (cluster-scoped)

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
   name: dev-team-cluster-read
rules:
   - apiGroups: [""]
   resources: ["nodes", "persistentvolumes"]
   verbs: ["get", "list", "watch"]
```

4. Create a ClusterRoleBinding for the Group

apiVersion: rbac.authorization.k8s.io/v1 kind: ClusterRoleBinding metadata: name: dev-team-clusterrolebinding subjects: - kind: Group # Must match your name: dev-team auth provider's group apiGroup: rbac.authorization.k8s.io roleRef: kind: ClusterRole name: dev-team-cluster-read apiGroup: rbac.authorization.k8s.io

Scenario:

You have a custom application running in a Pod that needs to **list and get**ConfigMaps and Secrets from the namespace it's deployed in. You don't want to give it too many permissions, just exactly what it needs. So, you'll create a dedicated ServiceAccount and use RBAC to allow access only to those resources.

Goal:

- Pod uses a specific **ServiceAccount**.
- That ServiceAccount has permissions to:
 - o get, list, and watch ConfigMaps and Secrets.
- Scoped only to the app-namespace.

★ Step-by-Step Setup

1. Create the ServiceAccount

apiVersion: v1

kind: ServiceAccount

metadata:

name: app-reader-sa

namespace: app-namespace

? This SA will be used by your application Pod.

2. Create the Role (namespace-scoped)

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
   name: read-secrets-configmaps
   namespace: app-namespace
rules:
   - apiGroups: [""]
   resources: ["secrets", "configmaps"]
   verbs: ["get", "list", "watch"]
```

This defines what the ServiceAccount is allowed to do in the namespace.

3. Bind the Role to the ServiceAccount

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: bind-sa-to-role
  namespace: app-namespace
subjects:
  - kind: ServiceAccount
    name: app-reader-sa
    namespace: app-namespace
roleRef:
  kind: Role
  name: read-secrets-configmaps
```

This links the permissions to the ServiceAccount.

apiGroup: rbac.authorization.k8s.io

4. Use the ServiceAccount in your Pod/Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: config-reader
  namespace: app-namespace
spec:
  replicas: 1
  selector:
    matchLabels:
      app: config-reader
  template:
    metadata:
      labels:
        app: config-reader
    spec:
      serviceAccountName: app-reader-sa
      containers:
        - name: app
```

image: my-app:latest

Now your app will only be able to read Secrets and ConfigMaps – nothing

else.

When a Kubernetes cluster is created, some default Roles, RoleBindings,

ClusterRoles, ClusterRoleBindings, and ServiceAccounts are automatically

created by Kubernetes to ensure basic functionality and security of the control

plane components and workloads.

* 1. Default ServiceAccounts

When you create a namespace (even default), Kubernetes automatically

creates a ServiceAccount named default.

Example:

apiVersion: v1

kind: ServiceAccount

metadata:

name: default

namespace: default

Purpose:

Every Pod without an explicitly defined serviceAccountName automatically uses this default ServiceAccount.

It typically has no permissions unless explicitly given via RBAC.

O

2. Default Roles & RoleBindings (Namespace-scoped)

By default, Kubernetes does not create many namespace-scoped Roles unless an addon or component (like the Dashboard) requires it.

But you can create your own Role and RoleBinding scoped to a namespace.

These are not created unless some specific component or workload installs them.

3. Default ClusterRoles

Kubernetes provides several **predefined ClusterRoles** to help manage access and permissions.

Examples of built-in ClusterRoles:

- cluster-admin Full control over the entire cluster
- admin Admin access within a namespace
- edit Can edit most resources in a namespace (not RBAC)

- view Read-only access in a namespace
- system:discovery For unauthenticated users to perform discovery (like listing API versions)
- system:node Used internally by kubelets
- system:controller:* For internal controllers like deployment-controller, job-controller, etc.

These roles are defined **cluster-wide**, and you can use them for your users, groups, or ServiceAccounts via **ClusterRoleBindings** or **RoleBindings**.

ClusterRoleBindings are automatically created for Kubernetes components to function properly.

Examples of default ClusterRoleBindings:

- cluster-admin Grants cluster-admin role to system:masters group (used by admins)
- system:node Binds system:node role to system:nodes group (all kubelets)

- system:public-info-viewer Lets unauthenticated users access basic cluster info
- system:controller:... Binds appropriate roles to internal controllers
- Represented the Properties of the Properties of
 - kubelets accessing the API server
 - Controllers managing deployments
 - Admin users performing cluster operations

Pro Tip:

You can **list** default roles and bindings in your cluster with:

kubectl get clusterrole
kubectl get clusterrolebinding
kubectl get role -A
kubectl get rolebinding -A
kubectl get serviceaccount -A

Real-Time RBAC Scenario: Dev & QA Access

Management in a Kubernetes Cluster

Use Case:

You are managing a Kubernetes cluster for a company with multiple teams:

- Dev Team: Should have full access to all resources only in the dev namespace.
- QA Team: Should have read-only access to Pods and Services in the qa namespace.
- A CI/CD tool (like Jenkins): Running in cicd namespace, needs permission to deploy apps in both dev and qa namespaces.

To implement this securely, you use **RBAC** with a combination of:

- Roles / ClusterRoles
- RoleBindings / ClusterRoleBindings
- Groups / Users / ServiceAccounts

Step-by-Step Implementation:

1 Dev Team Setup – Full Access in dev Namespace

Create a Role:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
   name: dev-team-role
   namespace: dev
rules:
- apiGroups: ["", "apps", "batch"]
   resources: ["pods", "services", "deployments", "jobs"]
   verbs: ["get", "list", "create", "delete", "update", "patch"]
```

☑ Bind Role to the Dev Group:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
```

name: dev-team-binding
namespace: dev
subjects:
- kind: Group
name: dev-team
apiGroup: rbac.authorization.k8s.io
roleRef:
 kind: Role
name: dev-team-role
apiGroup: rbac.authorization.k8s.io

This grants full permissions to any user in the dev-team group within the devnamespace.

2QA Team Setup – Read-Only in qa Namespace

Create a Role:

apiVersion: rbac.authorization.k8s.io/v1

kind: Role

```
metadata:
  name: qa-team-readonly
  namespace: qa
rules:
- apiGroups: [""]
  resources: ["pods", "services"]
  verbs: ["get", "list", "watch"]
Maind Role to QA Group:
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: qa-team-binding
  namespace: qa
subjects:
- kind: Group
  name: qa-team
  apiGroup: rbac.authorization.k8s.io
```

```
roleRef:
  kind: Role
  name: qa-team-readonly
  apiGroup: rbac.authorization.k8s.io
A This ensures the QA team can only view resources in the qa namespace.
3 CI/CD ServiceAccount Setup – Cluster-Wide Deployment Access
Create a ClusterRole:
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: cicd-deployer
rules:
- apiGroups: ["apps", ""]
  resources: ["deployments", "pods", "services"]
```

verbs: ["create", "update", "delete", "get", "list"]

✓ Create a ServiceAccount in cicd namespace:

apiVersion: v1 kind: ServiceAccount metadata: name: jenkins-deployer namespace: cicd **✓** Bind ClusterRole to ServiceAccount: apiVersion: rbac.authorization.k8s.io/v1 kind: ClusterRoleBinding metadata: name: cicd-deployer-binding subjects: - kind: ServiceAccount name: jenkins-deployer namespace: cicd roleRef:

kind: ClusterRole

name: cicd-deployer

apiGroup: rbac.authorization.k8s.io

This allows Jenkins (via the jenkins-deployer ServiceAccount) to deploy apps in any namespace across the cluster.

Summary of RBAC Entities Used:

Entity	Kind	Scope	Description
dev-team-rol e	Role	Namespa ce	Full access to Dev team
qa-team-read only	Role	Namespa ce	Read-only for QA team
cicd-deploye r	ClusterRole	Cluster	Full deploy access
RoleBinding	RoleBinding	Namespa ce	Bind Role to user/group
ClusterRoleB inding	ClusterRoleBindi ng	Cluster	Bind ClusterRole to ServiceAccount

Real-World Benefits of This Setup:

- Reast privilege access no team gets more permissions than required
- Clear separation of environments (dev, qa, prod)
- Enables automation tools (like Jenkins) to safely deploy apps
- Easy to audit and manage access per team/project

Objective

Set up **three different ServiceAccounts** with distinct permissions in a Kubernetes cluster, generate tokens, create dedicated kubeconfig files, and verify access from separate VMs/EC2s.

1. Create ServiceAccounts

kubectl create serviceaccount admin-sa -n default
kubectl create serviceaccount general-sa -n default
kubectl create serviceaccount others-sa -n default

2. Create ClusterRoles & ClusterRoleBindings

- A. Admin ServiceAccount Full Access
- ClusterRole:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: admin-cr
rules:
- apiGroups: ["*"]
  resources: ["*"]
  verbs: ["*"]
ClusterRoleBinding:
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: admin-crb
subjects:
- kind: ServiceAccount
  name: admin-sa
  namespace: default
roleRef:
```

```
kind: ClusterRole
name: admin-cr
apiGroup: rbac.authorization.k8s.io
```

☑ B. General ServiceAccount – Read-Only for Core Resources

ClusterRole:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
   name: general-cr
rules:
- apiGroups: [""]
   resources: ["pods", "services", "configmaps", "secrets"]
   verbs: ["get", "list", "watch"]
- apiGroups: ["apps"]
   resources: ["deployments"]
   verbs: ["get", "list", "watch"]
```

ClusterRoleBinding:

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
 name: general-crb
subjects:
- kind: ServiceAccount
 name: general-sa
 namespace: default
roleRef:
 kind: ClusterRole
 name: general-cr
 apiGroup: rbac.authorization.k8s.io

✓ C. Others ServiceAccount – Namespace-Level Read Access

ClusterRole:

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole

metadata:

```
name: others-cr
rules:
- apiGroups: [""]
  resources: ["namespaces"]
  verbs: ["get", "list", "watch"]
ClusterRoleBinding:
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: others-crb
subjects:
- kind: ServiceAccount
  name: others-sa
  namespace: default
roleRef:
  kind: ClusterRole
  name: others-cr
  apiGroup: rbac.authorization.k8s.io
```

3. Generate Tokens for ServiceAccounts

kubectl create token admin-sa -n default > admin.token
kubectl create token general-sa -n default > general.token
kubectl create token others-sa -n default > others.token

Save these tokens securely — we'll use them in kubeconfig files.

📜 4. Get Cluster Info for Kubeconfig

```
CLUSTER_NAME=$(kubectl config view --minify -o
jsonpath='{.clusters[0].name}')

CLUSTER_SERVER=$(kubectl config view --minify -o
jsonpath='{.clusters[0].cluster.server}')

CA_DATA=$(kubectl config view --minify -o
jsonpath='{.clusters[0].cluster.certificate-authority-data}')
```

5. Create Kubeconfig Files

Admin

```
cat <<EOF > kubeconfig-admin.yaml
apiVersion: v1
```

```
kind: Config
clusters:
- cluster:
    certificate-authority-data: $CA_DATA
    server: $CLUSTER_SERVER
  name: $CLUSTER_NAME
contexts:
- context:
    cluster: $CLUSTER_NAME
    user: admin
  name: admin-context
current-context: admin-context
users:
- name: admin
  user:
    token: $(cat admin.token)
EOF
```

General

```
cat <<EOF > kubeconfig-general.yaml
apiVersion: v1
kind: Config
clusters:
- cluster:
    certificate-authority-data: $CA_DATA
    server: $CLUSTER_SERVER
  name: $CLUSTER_NAME
contexts:
- context:
    cluster: $CLUSTER_NAME
    user: general
  name: general-context
current-context: general-context
users:
- name: general
  user:
```

```
token: $(cat general.token)
```

EOF

Others

```
cat <<EOF > kubeconfig-others.yaml
apiVersion: v1
kind: Config
clusters:
- cluster:
    certificate-authority-data: $CA_DATA
    server: $CLUSTER_SERVER
  name: $CLUSTER_NAME
contexts:
- context:
    cluster: $CLUSTER_NAME
    user: others
  name: others-context
current-context: others-context
```

```
users:
- name: others
user:
token: $(cat others.token)

EOF

6. Use from New EC2/VM
```

- - 1. Install kubectl
 - 2. Paste respective kubeconfig (kubeconfig-admin.yaml, etc.)

Export:

```
export KUBECONFIG=~/kubeconfig-admin.yaml # or
general/others
```

7. Verify Access in VM

Admin:

kubectl get pods -A

```
kubectl create ns test-ns
```

kubectl delete pod xyz -n dev

General:

kubectl get pods -n default

kubectl get configmaps

kubectl create ns test-ns \times (should be forbidden)

Others:

kubectl get namespaces

kubectl get pods -n default \times (should be forbidden)

What This RBAC Scenario Does

The scenario sets up three types of access levels in a Kubernetes cluster using ServiceAccounts, ClusterRoles, ClusterRoleBindings, and token-based authentication. Here's what it achieves:

You're assigning specific permissions to specific ServiceAccounts so that:

- Admin users can do anything in the cluster (create, delete, modify any resource).
- General users can only view important resources (pods, services, configmaps, secrets, deployments).
- Other users can **only view namespaces**, and nothing else.

This ensures **least privilege access** and prevents misuse or accidental modification of critical cluster resources.

🤦 2. Simulates Real-World Access Scenarios

Just like in a real company setup:

Admin team manages the entire cluster.

- **Dev or QA team** only needs read-only access for debugging.
- Auditors or compliance teams just need to verify what namespaces exist.

This kind of separation is **very common in production environments** and supports real-world **team-based access models**.

⊗ 3. Token-Based Authentication via ServiceAccounts

Instead of using kubeconfig generated for IAM or user credentials, you're using **ServiceAccount tokens** for authentication, which is:

- Lightweight
- Secure
- Works well for automation, bots, and internal tools

4. Custom Kubeconfig Files for External Access

You're generating **separate kubeconfig files** for each ServiceAccount and using them from **external EC2s or VMs**, simulating how:

- Devs connect from their own machines
- Bots or monitoring tools connect to the cluster

• Audit or Ops teams securely connect without sharing credentials

☑ 5. Complete Setup for Real Use Cases

This is a full-cycle implementation covering:

- Creation of ServiceAccounts
- Defining permissions (ClusterRole)
- Binding permissions (ClusterRoleBinding)
- Token generation
- Kubeconfig creation
- External machine access
- Access verification

It's like a **production-grade**, **best-practice example** of implementing secure and structured access to a Kubernetes cluster.

Why This Scenario Matters

- **General Security**: Avoid giving everyone cluster-admin access
- **Team-based access**: Dev, QA, Admin, Auditor all need different levels
- **Compliance**: Follows principle of least privilege
- Automation ready: Works great for bots and CI/CD tools
- Real-world experience: Prepares you for interviews, real jobs, and handling production workloads

@ What's Different in EKS?

Unlike KOPS, **EKS uses AWS IAM for authentication**, and **RBAC for authorization**. That means:

- You cannot directly use ServiceAccount tokens for human users to log in via kubectl.
- Instead, you authenticate using IAM identities (users or roles).
- Then, you map those IAM identities to Kubernetes users/groups using the aws-auth ConfigMap.

So, you achieve the **same separation of access and RBAC policies** — just with **IAM + RBAC combo** instead of pure ServiceAccount tokens.

Real-Time Equivalent Scenario in EKS

Let's convert your 3-service-account scenario into EKS-style:

1. IAM Users Setup

Create 3 IAM users in AWS:

- eks-admin
- eks-general
- eks-others

Generate access keys for each user.

• 2. Update aws-auth ConfigMap in EKS

Use the following command to edit:

kubectl edit configmap aws-auth -n kube-system

Add each user with appropriate groups:

```
mapUsers: |
  - userarn: arn:aws:iam::<account-id>:user/eks-admin
    username: eks-admin
    groups:
      - system:masters # Full access
  - userarn: arn:aws:iam::<account-id>:user/eks-general
    username: eks-general
    groups:
      - view-group
  - userarn: arn:aws:iam::<account-id>:user/eks-others
    username: eks-others
    groups:
      - ns-read-group
```

- 3. Create Corresponding ClusterRoles and ClusterRoleBindings
- ClusterRole for view-group

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: general-read-access
rules:
- apiGroups: [""]
  resources: ["pods", "services", "configmaps", "secrets"]
  verbs: ["get", "list", "watch"]
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["get", "list", "watch"]
ClusterRoleBinding:
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: general-read-binding
roleRef:
  kind: ClusterRole
  name: general-read-access
```

```
apiGroup: rbac.authorization.k8s.io
subjects:
- kind: Group
  name: view-group
  apiGroup: rbac.authorization.k8s.io
ClusterRole for ns-read-group
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: namespace-reader
rules:
- apiGroups: [""]
  resources: ["namespaces"]
  verbs: ["get", "list", "watch"]
ClusterRoleBinding:
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
```

```
name: ns-read-binding

roleRef:
    kind: ClusterRole
    name: namespace-reader
    apiGroup: rbac.authorization.k8s.io

subjects:
- kind: Group
    name: ns-read-group
    apiGroup: rbac.authorization.k8s.io
```

- You do **not** use tokens and kubeconfig manually here.
- Each IAM user runs aws eks update-kubeconfig with their credentials, and gets access as per RBAC.

Verifying Access from New EC2s/VMs

- On each EC2, configure credentials for respective IAM user using AWS CLI.
- Run:

```
aws eks update-kubeconfig --region <region> --name
<cluster-name>
```

Now test access:

- eks-admin → Full access
- eks-general → Only view workloads
- eks-others → Can only list namespace

How is RBAC Useful in Kubernetes?

RBAC (Role-Based Access Control) is one of the most powerful security features in Kubernetes. It allows you to control "who can do what" in your cluster, ensuring that only the right people or services have access to the right resources.

Here's how RBAC helps:

1. Enhanced Security

RBAC lets you follow the **Principle of Least Privilege (PoLP)** – meaning each user/service only gets **the minimal permissions needed** to do their job.

Example:

Your developer only needs to view deployments, not delete them. RBAC ensures that!

2. Team-Based Access Control

With RBAC, you can **define roles for different teams** — like developers, testers, admins, etc.

V Devs: Can view and edit resources in the dev namespace

Testers: Can only view logs and pods

Admins: Full access to everything

3. Centralized and Declarative Management

You manage access using YAML files, so it's **easy to review, audit, and version control** access policies — just like your code!

In EKS, RBAC works seamlessly with **AWS IAM**, so you can bind IAM users and roles to Kubernetes permissions — best of both worlds!

★ 5. Supports Multiple Subject Types

RBAC can be applied to:

- Users
- **Q** Groups
- X ServiceAccounts

So whether it's a human, a bot, or a CI/CD tool — you've got control!

4 6. Granular Access Control

You can define permissions for:

- Specific resources (Pods, Services, Secrets, etc.)
- Specific verbs (get, list, create, delete, etc.)
- Specific namespaces or the entire cluster

Want a service to only get and list pods in one namespace? RBAC can do that.

7. Prevents Costly Mistakes

RBAC protects your cluster from **accidental deletions or misconfigurations** by unauthorized users.

Imagine a junior developer accidentally deletes a production pod. RBAC can prevent that.

11 8. Auditing & Compliance

RBAC policies make it easier to **audit who has access to what**, which is super useful for:

- Security audits
- Compliance (HIPAA, GDPR, SOC 2)

How to Audit and Debug RBAC Issues in Kubernetes

When something doesn't work — like a user or service account can't access a resource — the problem is **often RBAC-related**. Here's how to **quickly identify and fix** those issues:

✓ 1. Use kubectl auth can-i – Your Best Friend

This command tells you **whether a user or service account** can perform a specific action.

Syntax:

kubectl auth can-i <verb> <resource> --as <user> -n
<namespace>

Example:

kubectl auth can-i get pods --as
system:serviceaccount:default:read-only-sa -n default

Use it to simulate access from any identity and quickly validate if the permissions are working.

2. Describe RoleBindings & ClusterRoleBindings

To verify which Roles/ClusterRoles are bound to a subject:

kubectl describe rolebinding <name> -n <namespace>
kubectl describe clusterrolebinding <name>

Check:

- **Subjects**: Are you binding the right user/group/service account?
- Roles: Do they have the correct rules?

✓ 3. View Roles & ClusterRoles

Sometimes a Role may not have the expected permissions. Check the details like this:

kubectl get role <name> -n <namespace> -o yaml
kubectl get clusterrole <name> -o yaml

✓ 4. Inspect ServiceAccounts

If you're using ServiceAccounts, make sure you're referencing them correctly:

kubectl get serviceaccount <name> -n <namespace> -o yaml

Confirm:

- Tokens are mounted correctly
- Name and namespace match the binding

✓ 5. Check API Server Logs (Advanced)

If you have access to the control plane (like in Kops), you can check audit logs:

Audit log location (self-managed cluster):

/var/log/kubernetes/audit.log

- Q Look for:
 - Denied requests
 - Forbidden errors ("reason": "Forbidden")

This is not possible in EKS unless CloudTrail or CloudWatch Audit Logging is enabled.

6. Enable Audit Logs in EKS (Optional for EKS)

If you're on **Amazon EKS**, enable Kubernetes audit logs in CloudWatch:

Go to:

EKS > Cluster > Logging > Enable Audit Logs

Use it to monitor all API requests and RBAC denials!

- **7.** Common Issues to Watch For
- NoleBinding in wrong namespace
- Fix: Ensure RoleBinding is in the same namespace as Role and subject
- **Solution** Using RoleBinding instead of ClusterRoleBinding
- Fix: Use ClusterRoleBinding if you're working across namespaces or want cluster-wide access
- Name or namespace
- Fix: Double-check both name and namespace when binding

🙌 Thank You for Being a Part of This RBAC Journey!

As we come to the end of this deep dive into **Kubernetes RBAC (Role-Based Access Control)**, I just want to take a moment to express my heartfelt gratitude to everyone who supported, engaged, and followed along.

Whether you're a beginner just getting started with Kubernetes or an experienced engineer refining your security practices — I hope this guide gave you the clarity and confidence to implement RBAC effectively and securely in your clusters.

What we've achieved together in this doc:

- We broke down RBAC from the ground up theory to practice.
- Explored real-world scenarios using Users, Groups, and ServiceAccounts.
- Set up complete role-based environments with scoped access control.
- Learned how to debug and audit permissions to stay in control.
- And finally, built a secure and scalable access management layer tailored for real Kubernetes clusters (Kops & EKS).

What's Next?

The journey doesn't stop here!

There's **so much more in the Kubernetes world** that I can't wait to explore and share with you.

From **Network Policies**, **Pod Security**, **Helm**, **GitOps**, to **CI/CD in Kubernetes** — I'm just getting started.

So if you've enjoyed this doc:

- **Let's connect** I'd love to hear your feedback and suggestions.
- Share it with others who might find this useful.
- Follow me for more content daily Kubernetes learnings, hands-on projects, AWS tasks, and DevOps deep dives are on the way.

* Thank You Again

Your support truly means the world to me — every comment, like, message, and reshare encourages me to keep going and keep learning.

Let's continue growing, building, and helping each other in this amazing DevOps & Cloud community.

With tons of gratitude,

~ Fenil Gajjar

"Keep learning. Keep building. And always stay curious. 🚀"